

Problem statement:To predict How Best the data fits

1. Data collection

In [33]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import preprocessing,svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

In [34]:

```
df=pd.read_csv(r"C:\Users\jyothi reddy\Downloads\insurance.csv")
df
```

Out[34]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

2.Data cleaning and Preprocessing

#Exploratory data anlysis

In [35]:

```
df.head()
```

Out[35]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

In [36]:

```
df.tail()
```

Out[36]:

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

In [37]:

```
df.shape
```

Out[37]:

(1338, 7)

In [38]:

```
df.describe
```

Out[38]:

<bound method NDFrame.describe of  
0 19 female 27.900 0 yes southwest 16884.92400  
1 18 male 33.770 1 no southeast 1725.55230  
2 28 male 33.000 3 no southeast 4449.46200  
3 33 male 22.705 0 no northwest 21984.47061  
4 32 male 28.880 0 no northwest 3866.85520  
... ... ... ... ... ... ...  
1333 50 male 30.970 3 no northwest 10600.54830  
1334 18 female 31.920 0 no northeast 2205.98080  
1335 18 female 36.850 0 no southeast 1629.83350  
1336 21 female 25.800 0 no southwest 2007.94500  
1337 61 female 29.070 0 yes northwest 29141.36030  
  
[1338 rows x 7 columns]>

In [39]:

```
df.info()
```

<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1338 entries, 0 to 1337  
Data columns (total 7 columns):  
# Column Non-Null Count Dtype  
--- ---  
0 age 1338 non-null int64  
1 sex 1338 non-null object  
2 bmi 1338 non-null float64  
3 children 1338 non-null int64  
4 smoker 1338 non-null object  
5 region 1338 non-null object  
6 charges 1338 non-null float64  
dtypes: float64(2), int64(2), object(3)  
memory usage: 73.3+ KB

In [40]:

```
df.isnull().any()
```

Out[40]:

age False  
sex False  
bmi False  
children False  
smoker False  
region False  
charges False  
dtype: bool

In [41]:

```
df.isna().sum()
```

Out[41]:

age 0  
sex 0  
bmi 0  
children 0  
smoker 0  
region 0  
charges 0  
dtype: int64

In [42]:

```
df['region'].value_counts()
```

Out[42]:

```
region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64
```

In [43]:

```
convert={"sex":{"female":1,"male":0}}
df=df.replace(convert)
df
```

Out[43]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	yes	southwest	16884.92400
1	18	0	33.770	1	no	southeast	1725.55230
2	28	0	33.000	3	no	southeast	4449.46200
3	33	0	22.705	0	no	northwest	21984.47061
4	32	0	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	0	30.970	3	no	northwest	10600.54830
1334	18	1	31.920	0	no	northeast	2205.98080
1335	18	1	36.850	0	no	southeast	1629.83350
1336	21	1	25.800	0	no	southwest	2007.94500

In [44]:

```
convert={"smoker":{"yes":1, "no":0}}
df=df.replace(convert)
df
```

Out[44]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	southwest	16884.92400
1	18	0	33.770	1	0	southeast	1725.55230
2	28	0	33.000	3	0	southeast	4449.46200
3	33	0	22.705	0	0	northwest	21984.47061
4	32	0	28.880	0	0	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	0	30.970	3	0	northwest	10600.54830
1334	18	1	31.920	0	0	northeast	2205.98080
1335	18	1	36.850	0	0	southeast	1629.83350
1336	21	1	25.800	0	0	southwest	2007.94500

In [45]:

```
convert={"region":{"southwest":1,"southeast":2,"northwest":3,"northeast":4}}
df=df.replace(convert)
df
```

Out[45]:


	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	1	16884.92400
1	18	0	33.770	1	0	2	1725.55230
2	28	0	33.000	3	0	2	4449.46200
3	33	0	22.705	0	0	3	21984.47061
4	32	0	28.880	0	0	3	3866.85520
...	...	...	...	...	...	...	...
1333	50	0	30.970	3	0	3	10600.54830
1334	18	1	31.920	0	0	4	2205.98080
1335	18	1	36.850	0	0	2	1629.83350
1336	21	1	25.800	0	0	1	2007.94500
1337	61	1	29.070	0	1	3	29141.36030

1338 rows × 7 columns

### 3.Data Visualization

In [46]:

```
sns.lmplot(x='bmi',y='charges',order=2,data=df,ci=None)
plt.show()
```



In [47]:

```
x=np.array(df['bmi']).reshape(-1,1)
y=x*np.array(df['charges']).reshape(-1,1)
```


In [48]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=0)
lr=LinearRegression()
lr.fit(x_train,y_train)
print(lr.score(x_test,y_test))
```

1.0

In [49]:


```
y_pred=lr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



## working with subset of data

In [50]:

```
df700=df[:][:700]
sns.lmplot(x='bmi',y='charges',order=2,ci=None,data=df700)
plt.show()
```



In [51]:

```
df700.fillna(method='ffill',inplace=True)
```

In [52]:

```
x=np.array(df700["bmi"]).reshape(-1,1)
y=np.array(df700['charges']).reshape(-1,1)
```

In [53]:

```
df700.dropna(inplace=True)
```


In [54]:

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
lr=LinearRegression()
lr.fit(x_train,y_train)
print(lr.score(x_test,y_test))
```

```
-0.1630229146000015
```

In [55]:

```
y_pred=lr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



## Evaluation of model

In [56]:

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
```

In [57]:

```
lr=LinearRegression()
lr.fit(x_train,y_train)
y_pred=lr.predict(x_test)
r2=r2_score(y_test,y_pred)
print(r2)
```

```
-0.1630229146000015
```

## Ridge Regression

In [58]:

```
from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

In [59]:

```
plt.figure(figsize=(10,10))
sns.heatmap(df700.corr(),annot=True)
plt.show()
```



In [60]:

```
features=df.columns[0:1]
target=df.columns[-1]
```

In [61]:

```
x=df[features].values
y=df[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state=1)
print("The dimension of X_train is {}".format(x_train.shape))
print("The dimension of X_test is {}".format(x_test.shape))
```

The dimension of X\_train is (936, 1)  
The dimension of X\_test is (402, 1)

In [62]:

```
lr = LinearRegression()  
#Fit model  
lr.fit(x_train, y_train)  
#predict  
actual = y_test  
train_score_lr = lr.score(x_train, y_train)  
test_score_lr = lr.score(x_test, y_test)  
print("\nLinear Regression Model:\n")  
print("The train score for lr model is {}".format(train_score_lr))  
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.0910963973805714

The test score for lr model is 0.08490473916580776

In [63]:

```
ridgeReg = Ridge(alpha=10)  
ridgeReg.fit(x_train,y_train)  
#train and test score for ridge regression  
train_score_ridge = ridgeReg.score(x_train, y_train)  
test_score_ridge = ridgeReg.score(x_test, y_test)  
print("\nRidge Model:\n")  
print("The train score for ridge model is {}".format(train_score_ridge))  
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.09109639711159634

The test score for ridge model is 0.08490538609860176

In [64]:

```
plt.figure(figsize=(10,10))
```

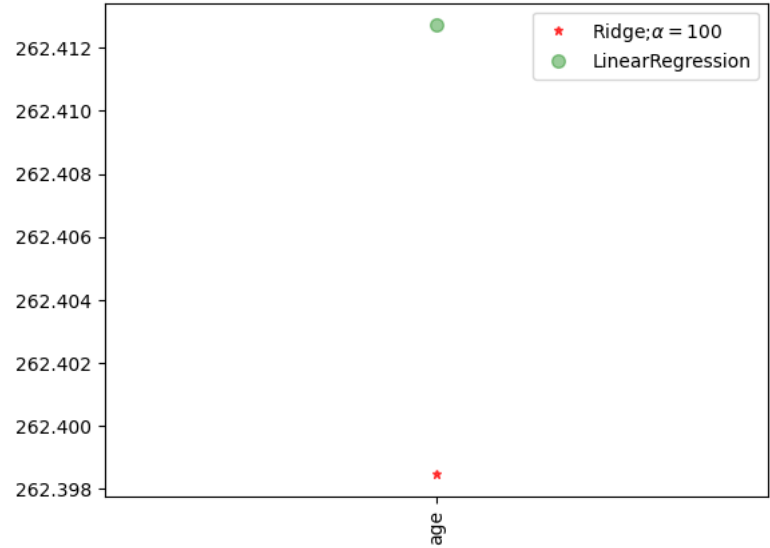
Out[64]:

<Figure size 1000x1000 with 0 Axes>

<Figure size 1000x1000 with 0 Axes>

In [66]:

```
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker="*",markersize=5,color='red',label=r'Ridge;$\alpha=100$')  
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker="o",markersize=7,color='green',label='LinearRegression')  
plt.xticks(rotation=90)  
plt.legend()  
plt.show()
```



## Lasso Regression



In [67]:

```
lasso= Lasso(alpha=10)
lasso.fit(x_train,y_train)
#train and test score for ridge regression
train_score_ls = lasso.score(x_train, y_train)
test_score_ls= lasso.score(x_test, y_test)
print("\nRidge Model:\n")
print("The train score for lasso model is {}".format(train_score_ls))
print("The test score for lasso model is {}".format(test_score_ls))
```

Ridge Model:

The train score for lasso model is 0.09109639395809055  
 The test score for lasso model is 0.08490704421828055

In [68]:

```
plt.figure(figsize=(10,10))
```


Out[68]:

&lt;Figure size 1000x1000 with 0 Axes&gt;

&lt;Figure size 1000x1000 with 0 Axes&gt;

In [69]:

```
pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")
plt.show()
```



In [70]:

```
from sklearn.linear_model import LassoCV
```

In [71]:

```
#using the linear cv model
from sklearn.linear_model import RidgeCV
#cross validation
ridge_cv=RidgeCV(alphas =[0.0001,0.001,0.01,0.1,1,10]).fit(x_train,y_train)
#score
print(ridge_cv.score(x_train,y_train))
print(ridge_cv.score(x_test,y_test))
```

0.09109639711159612  
 0.08490538609884779


In [72]:

```
#using the linear cv model
from sklearn.linear_model import LassoCV
#cross validation
lasso_cv=LassoCV(alphas =[0.0001,0.001,0.01,0.1,1,10]).fit(x_train,y_train)
#score
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

```
0.09109639395809055
0.08490704421828055
```

In [73]:

```
plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge; $\alpha=100$')
#add plot for lasso regression
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso; $\alpha = \text{grid}$')
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
#rotate axis
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
plt.show()
```



## ElasticNet Regression

In [74]:

```
from sklearn.linear_model import ElasticNet
```

In [75]:

```
el=ElasticNet()  
el.fit(x_train,y_train)  
print(el.coef_)  
print(el.intercept_)
```

[261.74450967]  
3115.083177426244

In [76]:

```
y_pred_elastic=el.predict(x_train)
```

In [77]:

```
mean_squared_error=np.mean((y_pred_elastic-y_train)**2)  
print(mean_squared_error)
```

135077142.70714515

In [78]:

```
el=ElasticNet()  
el.fit(x_train,y_train)  
print(el.score(x_train,y_train))
```

0.09109580670592365

## Logistic Regression

In [79]:

```
import numpy as np  
import pandas as pd  
from sklearn.linear_model import LogisticRegression  
from sklearn.preprocessing import StandardScaler
```

In [80]:

```
df=pd.read_csv(r"C:\Users\jyothi reddy\Downloads\insurance.csv")  
df
```

Out[80]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...	...	...	...	...	...	...	...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

In [81]:

```
df.shape
```

Out[81]:

(1338, 7)

In [82]:

```
pd.set_option('display.max_rows',1000000000)
pd.set_option('display.max_columns',1000000000)
pd.set_option('display.width',95)
```

In [83]:

```
print('This Dataset has %d rows and %d columns'%(df.shape))
```

This Dataset has 1338 rows and 7 columns

In [84]:

```
df.head()
```

Out[84]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

In [85]:

```
df.describe
```

42	41	male	21.780		1	no	southeast	6272.477200
43	37	female	30.800		2	no	southeast	6313.759000
44	38	male	37.050		1	no	northeast	6079.671500
45	55	male	37.300		0	no	southwest	20630.283510
46	18	female	38.665		2	no	northeast	3393.356350
47	28	female	34.770		0	no	northwest	3556.922300
48	60	female	24.530		0	no	southeast	12629.896700
49	36	male	35.200		1	yes	southeast	38709.176000
50	18	female	35.625		0	no	northeast	2211.130750
51	21	female	33.630		2	no	northwest	3579.828700
52	48	male	28.000		1	yes	southwest	23568.272000
53	36	male	34.430		0	yes	southeast	37742.575700
54	40	female	28.690		3	no	northwest	8059.679100
55	58	male	36.955		2	yes	northwest	47496.494450
56	58	female	31.825		2	no	northeast	13607.368750
57	18	male	31.680		2	yes	southeast	34303.167200
58	53	female	22.880		1	yes	southeast	23244.790200
59	34	female	37.335		2	no	northwest	5989.523650
60	43	male	27.360		3	no	northeast	8606.217400
61	25	male	33.660		4	no	southeast	4504.662400

In [86]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   age         1338 non-null   int64  
 1   sex         1338 non-null   object  
 2   bmi         1338 non-null   float64 
 3   children    1338 non-null   int64  
 4   smoker      1338 non-null   object  
 5   region      1338 non-null   object  
 6   charges     1338 non-null   float64 
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [87]:

```
df.isnull().sum()
```

Out[87]:

```
age      0
sex      0
bmi      0
children 0
smoker   0
region   0
charges  0
dtype: int64
```

In [88]:

```
convert={"smoker":{"yes":1,"no":0}}
df=df.replace(convert)
df
```

2	28	male	33.000	3	0	southeast	4449.462000
3	33	male	22.705	0	0	northwest	21984.470610
4	32	male	28.880	0	0	northwest	3866.855200
5	31	female	25.740	0	0	southeast	3756.621600
6	46	female	33.440	1	0	southeast	8240.589600
7	37	female	27.740	3	0	northwest	7281.505600
8	37	male	29.830	2	0	northeast	6406.410700
9	60	female	25.840	0	0	northwest	28923.136920
10	25	male	26.220	0	0	northeast	2721.320800
11	62	female	26.290	0	1	southeast	27808.725100
12	23	male	34.400	0	0	southwest	1826.843000
13	56	female	39.820	0	0	southeast	11090.717800
14	27	male	42.130	0	1	southeast	39611.757700

In [89]:

```
convert={"sex":{"female":1,"male":0}}
df=df.replace(convert)
df
```

Out[89]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	southwest	16884.924000
1	18	0	33.770	1	0	southeast	1725.552300
2	28	0	33.000	3	0	southeast	4449.462000
3	33	0	22.705	0	0	northwest	21984.470610
4	32	0	28.880	0	0	northwest	3866.855200
5	31	1	25.740	0	0	southeast	3756.621600
6	46	1	33.440	1	0	southeast	8240.589600
7	37	1	27.740	3	0	northwest	7281.505600
8	37	0	29.830	2	0	northeast	6406.410700
9	60	1	25.840	0	0	northwest	28923.136920

In [90]:

```
convert={"region":{"southeast":1,"southwest":2,"northeast":3,"northwest":4}}
df=df.replace(convert)
df
```

1	18	0	33.770	1	0	1	1725.552300
2	28	0	33.000	3	0	1	4449.462000
3	33	0	22.705	0	0	4	21984.470610
4	32	0	28.880	0	0	4	3866.855200
5	31	1	25.740	0	0	1	3756.621600
6	46	1	33.440	1	0	1	8240.589600
7	37	1	27.740	3	0	4	7281.505600
8	37	0	29.830	2	0	3	6406.410700
9	60	1	25.840	0	0	4	28923.136920
10	25	0	26.220	0	0	3	2721.320800
11	62	1	26.290	0	1	1	27808.725100
12	23	0	34.400	0	0	2	1826.843000
13	56	1	39.820	0	0	1	11090.717800

In [91]:

```
features_matrix=df.iloc[:,0:4]
```

In [92]:

```
target_vector=df.iloc[:,-3]
```

In [93]:

```
print('The Feature Matrix has %d Rows and %d columns(s)%(features_matrix.shape))  
print('The Target Matrix has %d Rows and %d columns(s)%(np.array(target_vector).reshape(-1,1).shape))
```


The Feature Matrix has 1338 Rows and 4 columns(s)  
The Target Matrix has 1338 Rows and 1 columns(s)

In [94]:

```
import matplotlib.pyplot as plt  
import seaborn as sns
```

In [95]:

```
sns.barplot(x='smoker', y='age', data=df, color="mediumturquoise")  
plt.show()
```



In [96]:

```
features_matrix_standardized=StandardScaler().fit_transform(features_matrix)
```

In [97]:

```
algorithm=LogisticRegression(max_iter=10000)
```

In [98]:

```
Logistic_Regression_Model=algorithm.fit(features_matrix_standardized,target_vector)
```

In [99]:

```
observation=[[1,0,0.99539,-0.0588]]
```

In [101]:

```
predictions=Logistic_Regression_Model.predict(observation)  
print('The model predicted the observation to belong to class %s'%(predictions))
```

The model predicted the observation to belong to class [0]

In [102]:

```
print('The algorithm was trained to predict one of the two classes:%s'%(algorithm.classes_))
```

The algorithm was trained to predict one of the two classes:[0 1]

In [104]:

"The Model says the probability of the observation we passed belonging to class[0] %s" % "%(algorithm.predict\_proba(observation)[0][0]))"

The Model says the probability of the observation we passed belonging to class[0] 0.8057075871331396

In [107]:

Model says the probability of the observation we passed belonging to class['g'] Is %s" % "%(algorithm.predict\_proba(observation)[0][1]))"

The Model says the probability of the observation we passed belonging to class['g'] Is 0.19429241286686041

In [108]:

x=np.array(df['age']).reshape(-1,1)  
y=np.array(df['smoker']).reshape(-1,1)

In [109]:

x\_train,x\_test,y\_train,y\_test=train\_test\_split(x,y,test\_size=0.05)  
lo=LogisticRegression()  
lo.fit(x\_train,y\_train)  
print(lo.score(x\_test,y\_test))

0.7611940298507462

C:\Users\jyothi reddy\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().  
y = column\_or\_1d(y, warn=True)

## Decision Tree

In [110]:

import numpy as np  
import pandas as pd  
import seaborn as sns  
from sklearn.model\_selection import train\_test\_split  
from sklearn.tree import DecisionTreeClassifier

In [111]:

df=pd.read\_csv(r"C:\Users\jyothi reddy\Downloads\insurance.csv")  
df

4	32	male	28.880	0	no	northwest	3866.855200
5	31	female	25.740	0	no	southeast	3756.621600
6	46	female	33.440	1	no	southeast	8240.589600
7	37	female	27.740	3	no	northwest	7281.505600
8	37	male	29.830	2	no	northeast	6406.410700
9	60	female	25.840	0	no	northwest	28923.136920
10	25	male	26.220	0	no	northeast	2721.320800
11	62	female	26.290	0	yes	southeast	27808.725100
12	23	male	34.400	0	no	southwest	1826.843000
13	56	female	39.820	0	no	southeast	11090.717800
14	27	male	42.130	0	yes	southeast	39611.757700
15	19	male	24.600	1	no	southwest	1837.237000
16	52	female	30.780	1	no	northeast	10797.336200

In [112]:

df.shape

Out[112]:  
(1338, 7)

In [113]:

```
df.isnull().any()
```

Out[113]:

```
age      False
sex      False
bmi      False
children False
smoker   False
region   False
charges  False
dtype: bool
```

In [114]:

```
df['region'].value_counts()
```

Out[114]:

```
region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64
```

In [115]:

```
convert={"sex":{"female":1,"male":0}}
df=df.replace(convert)
df
```

20	60	1	36.005	0	no	northeast	13228.846950
21	30	1	32.400	1	no	southwest	4149.736000
22	18	0	34.100	0	no	southeast	1137.011000
23	34	1	31.920	1	yes	northeast	37701.876800
24	37	0	28.025	2	no	northwest	6203.901750
25	59	1	27.720	3	no	southeast	14001.133800
26	63	1	23.085	0	no	northeast	14451.835150
27	55	1	32.775	2	no	northwest	12268.632250
28	23	0	17.385	1	no	northwest	2775.192150
29	31	0	36.300	2	yes	southwest	38711.000000
30	22	0	35.600	0	yes	southwest	35585.576000
31	18	1	26.315	0	no	northeast	2198.189850
32	19	1	28.600	5	no	southwest	4687.797000
...	...	...	...	...	...	...	...

In [116]:

```
convert={"smoker":{"yes":1,"no":0}}
df=df.replace(convert)
df
```

19	30	0	35.300	0	1	southwest	36837.467000
20	60	1	36.005	0	0	northeast	13228.846950
21	30	1	32.400	1	0	southwest	4149.736000
22	18	0	34.100	0	0	southeast	1137.011000
23	34	1	31.920	1	1	northeast	37701.876800
24	37	0	28.025	2	0	northwest	6203.901750
25	59	1	27.720	3	0	southeast	14001.133800
26	63	1	23.085	0	0	northeast	14451.835150
27	55	1	32.775	2	0	northwest	12268.632250
28	23	0	17.385	1	0	northwest	2775.192150
29	31	0	36.300	2	1	southwest	38711.000000
30	22	0	35.600	0	1	southwest	35585.576000
31	18	1	26.315	0	0	northeast	2198.189850



In [117]:

```
x=["bmi","children"]
y=["Yes","No"]
all_inputs=df[x]
all_classes=df["sex"]
```

In [118]:

```
(x_train,x_test,y_train,y_test)=train_test_split(all_inputs,all_classes,test_size=0.03)
```

In [119]:

```
clf=DecisionTreeClassifier(random_state=0)
```

In [120]:

```
clf.fit(x_train,y_train)
```

Out[120]:

▼

DecisionTreeClassifier

DecisionTreeClassifier(random\_state=0)

In [121]:

```
score=clf.score(x_test,y_test)
print(score)
```

0.5853658536585366

## Random Forest

In [122]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt ,seaborn as sns
```

In [123]:

```
df=pd.read_csv(r"C:\Users\jyothi reddy\Downloads\insurance.csv")
df
```

34	28	male	36.400	1	yes	southwest	51194.559140
35	19	male	20.425	0	no	northwest	1625.433750
36	62	female	32.965	3	no	northwest	15612.193350
37	26	male	20.800	0	no	southwest	2302.300000
38	35	male	36.670	1	yes	northeast	39774.276300
39	60	male	39.900	0	yes	southwest	48173.361000
40	24	female	26.600	0	no	northeast	3046.062000
41	31	female	36.630	2	no	southeast	4949.758700
42	41	male	21.780	1	no	southeast	6272.477200
43	37	female	30.800	2	no	southeast	6313.759000
44	38	male	37.050	1	no	northeast	6079.671500
45	55	male	37.300	0	no	southwest	20630.283510
46	18	female	38.665	2	no	northeast	3393.356350

In [124]:

```
df.shape
```

Out[124]:

(1338, 7)

In [125]:

```
df['region'].value_counts()
```

Out[125]:

```
region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64
```

In [126]:

```
df['bmi'].value_counts()
```

```
25.800    7
32.775    7
27.645    7
32.110    7
38.060    7
25.460    7
30.590    7
27.360    7
24.320    7
34.800    7
27.500    6
19.950    6
29.920    6
30.115    6
26.600    6
30.200    6
35.530    6
33.630    6
28.595    6
27.100    6
```

In [127]:

```
m={"sex":{"female":1,"male":0}}
df=df.replace(m)
print(df)
```

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	yes	southwest	16884.924000
1	18	0	33.770	1	no	southeast	1725.552300
2	28	0	33.000	3	no	southeast	4449.462000
3	33	0	22.705	0	no	northwest	21984.470610
4	32	0	28.880	0	no	northwest	3866.855200
5	31	1	25.740	0	no	southeast	3756.621600
6	46	1	33.440	1	no	southeast	8240.589600
7	37	1	27.740	3	no	northwest	7281.505600
8	37	0	29.830	2	no	northeast	6406.410700
9	60	1	25.840	0	no	northwest	28923.136920
10	25	0	26.220	0	no	northeast	2721.320800
11	62	1	26.290	0	yes	southeast	27808.725100
12	23	0	34.400	0	no	southwest	1826.843000
13	56	1	39.820	0	no	southeast	11090.717800
14	27	0	42.130	0	yes	southeast	39611.757700
15	19	0	24.600	1	no	southwest	1837.237000
16	52	1	30.780	1	no	northeast	10797.336200
17	23	0	23.845	0	no	northeast	2395.171550

In [128]:

```
n={"smoker":{"yes":1,"no":0}}
df=df.replace(n)
print(df)
```

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	southwest	16884.924000
1	18	0	33.770	1	0	southeast	1725.552300
2	28	0	33.000	3	0	southeast	4449.462000
3	33	0	22.705	0	0	northwest	21984.470610
4	32	0	28.880	0	0	northwest	3866.855200
5	31	1	25.740	0	0	southeast	3756.621600
6	46	1	33.440	1	0	southeast	8240.589600
7	37	1	27.740	3	0	northwest	7281.505600
8	37	0	29.830	2	0	northeast	6406.410700
9	60	1	25.840	0	0	northwest	28923.136920
10	25	0	26.220	0	0	northeast	2721.320800
11	62	1	26.290	0	1	southeast	27808.725100
12	23	0	34.400	0	0	southwest	1826.843000
13	56	1	39.820	0	0	southeast	11090.717800
14	27	0	42.130	0	1	southeast	39611.757700
15	19	0	24.600	1	0	southwest	1837.237000
16	52	1	30.780	1	0	northeast	10797.336200
17	23	0	23.845	0	0	northeast	2395.171550

In [129]:

```
from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[129]:

RandomForestClassifier

RandomForestClassifier()

In [130]:

```
rf=RandomForestClassifier()
params={'max_depth':[2,3,5,20],
        'min_samples_leaf':[5,10,20,50,100,200],
        'n_estimators':[10,25,30,50,100,200]}
```

In [132]:

```
from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

Out[132]:

GridSearchCV

estimator: RandomForestClassifier

RandomForestClassifier

In [133]:

```
grid_search.best_score_
```

Out[133]:

0.5096254446536932


In [134]:

```
rf_best=grid_search.best_estimator_
print(rf_best)
```

RandomForestClassifier(max\_depth=2, min\_samples\_leaf=100, n\_estimators=10)

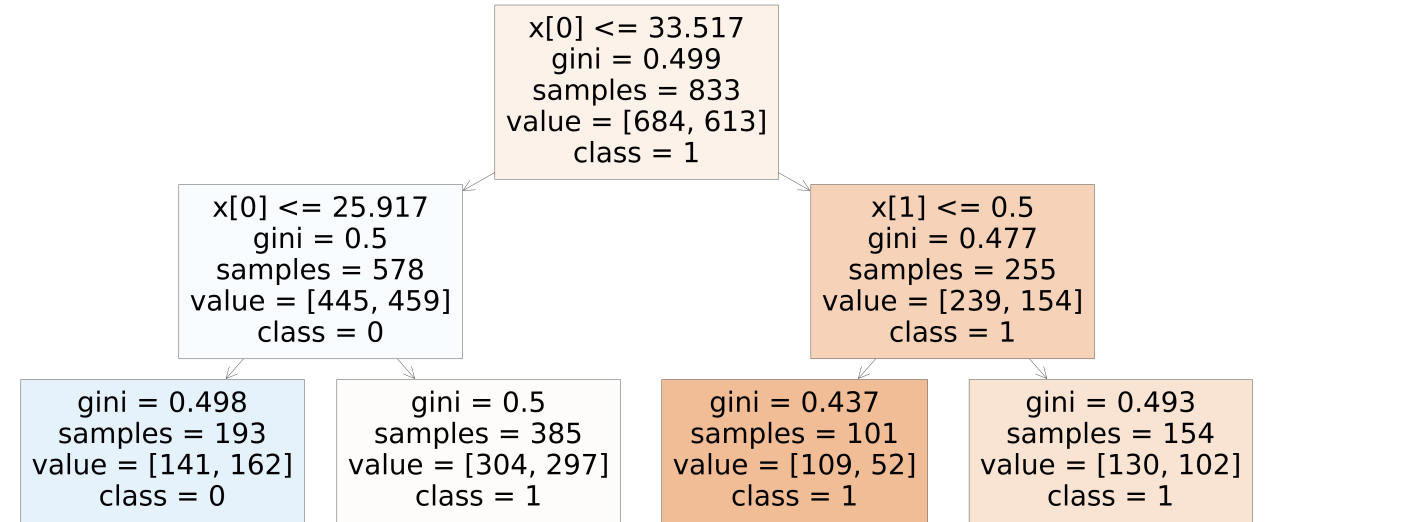
In [135]:

```
from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[4],class_names=['1','0'],filled=True);
```



In [136]:

```
from sklearn.tree import plot_tree
plt.figure(figsize=(70,30))
plot_tree(rf_best.estimators_[6],class_names=["1","0"],filled=True);
```



In [137]:

```
rf_best.feature_importances_
```

Out[137]:

```
array([0.8019108, 0.1980892])
```

In [140]:

```
rf=RandomForestClassifier(random_state=0)
```

In [141]:

```
rf.fit(x_train,y_train)
```

Out[141]:

```
RandomForestClassifier
RandomForestClassifier(random_state=0)
```

In [142]:

```
score=rf.score(x_test,y_test)
print(score)
```

```
0.5609756097560976
```

In [ ]: